

INTRODUCTION TO ARTIFICIAL INTELLIGENCE

1. THE AI PROBLEM

fundamental challenges and inherent difficulties face
when attempting to create truly intelligent artificial
system

Key Aspects

MAKING AI
UNDERSTAND WHATS
INTELLIGENCE &
KNOWLEDGE

Defining Intelligence for machines:

how do explain a machine what is intelligence
, because it is about many things like facts,
learning, reasoning, creativity, problem-solving,
social understanding, emotional intelligence and m
e.g. how a program understands and takes jokes
feel empathy.

How do you translate this into algo and data
structures

The common sense problem

X COMMON
SENSE

For humans there are many unaid intuitive
knowledge which are to be considered b
doing something, the COMMON SENSE
It's for hard to train AI with all of th

e.g. AI might learn from images that cars drive on road, but lacks common sense that car won't drive off a cliff (x self knowledge)

Impact on Env after an action

• The Frame Problem

when an AI performs an action or senses a change in its environment, it needs to understand what aspects of the world have changed and what's unchanged.

"About AI knowing what exactly changes and what doesn't in the environment after an action"
e.g. If a robot moves coffee mug from a table, it needs to know that the table's color didn't change, the room didn't rotate, and the gravity is still working and much more.

Impossible to cover all possible conditions

• The Qualification Problem

This difficulty is about fully specifying all the conditions, prerequisites, and exceptions for the rule or action to be valid

"About AI dealing with the fact that there are too many possible conditions to list for an action to work"

e.g. A robot instructed to open the door must implicitly assume the door isn't locked, isn't stuck, isn't painted onto the wall, isn't too heavy etc. Listing all these is too much and resulting in failure in unexpected situations

MISUSE AND X PRIVACY

The AI control problem
challenge of ensuring that advance AIs operate in a way that is beneficial, safe, and aligned with human values, ethics and intentions.
"Designing AI with goal alignment to prevent unintended and potentially harmful consequences."
e.g. using AI to hack on criminal activities
AI accidentally supporting harmful thoughts and encourage them
AI leaking sensitive data

JOB DISPLACEMENT
BIAS OUTCOMES
IF DATA BIASED

Ethical Dilemmas and Societal Impact
Concerns about AI bias (training data isn't always correct for all situations), privacy, job loss, no accountability for AI decisions and potential misuse
e.g. AI used for predicting criminal behaviour might inadvertently perpetuate existing social biases
if trained on historical biased data, leading unfair outcomes

THE UNDERLYING ASSUMPTION

"belief or principal that is taken for granted, often without explicit statement or proof"

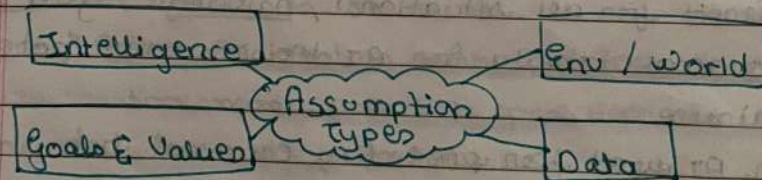
- forms unstated foundation for reasoning, actions or design of a system
- we (humans) subconsciously or consciously accept as true to make progress
- why is it critical?

for solving or acting upon a problem
we inherently assume few things about the world
AI will operate in, like amount, type of data
and behaviours and more.

These assumptions
shape entire design of the AI

- algorithm chosen
- type and amount of data collected
- how system learns
- how it performs and its measurement

If flawed everything goes wrong.



1. Assumption about Environment & World

"relates to characteristics of the operational environment for an AI agent"

- e.g.
1. AI agent designed for playing chess assumed a perfectly known, deterministic, discrete and static env (chess board and rules)
 2. Self-driving car AI operates in a high dynamic, continuous and partially observable environment. It must assume certain road rules, predictable pedestrian behaviour and general weather patterns etc

2. Assumption about Data

- ↳ Data Completeness: AI might assume its training data covers everything
- ↳ Data Accuracy: it often assumes data is free from errors etc
- ↳ Data Distribution: assume that data will be ~~some~~ similar to training data
- ↳ Data Bias: AI assumes training data is fair and true

3. Assumptions about Goals & Values

"what we ^{define} as success, optimal behaviour or ultimate AI objective"

- e.g. AI designed purely maximise profit for a company might implicitly assume that short term financial gain is sole objective. (so AI determines CROA itself)

4. Assumptions about Intelligence

"How intelligence works or how it can be computationally implemented"

e.g. Early AI assumed intelligence could be captured by explicit rules, logic and symbols
Modern AI assumes intelligence emerges from learning complex patterns in vast amount of data, without explicit rules and that sufficient data and computational power will lead to intelligent behaviour.

Impact of Unstated or Incorrect Assumptions

- 1. Inaccuracy
- 2. Underperformance
- 3. Ethical issues (bias)
- 4. Unintended consequences (when in business one e.g. 11)

AI TECHNIQUES

"Specific methods, algorithms and approaches that allow artificial intelligence systems to do tasks that traditionally required human intelligence"

Categories

1. Symbolic AI (Rule-Based Systems)

Focuses on representing human knowledge in symbolic, explicit form often using human defined rules.

e.g. expert system designed to diagnose car problems
so rules defined like

IF engine won't start AND battery dead
THEN check battery terminals

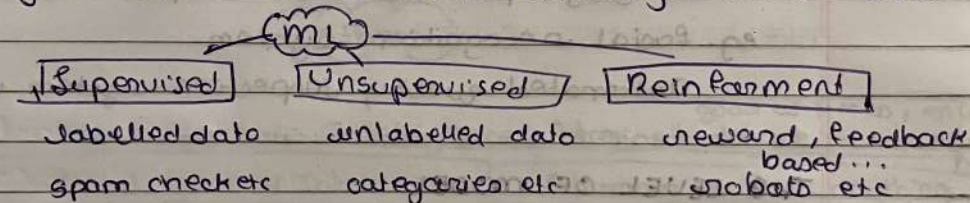
so basically uses collection of if-then

2. Machine Learning

- enables computers to learn from data without being explicitly programmed for everything.

- Finds patterns and make predictions or decisions

Give most of your ML knowledge here like



3. Search Algorithms

used to explore a set of possible solutions or paths to find a specific goal

e.g. GPS, road track, self driven cars etc

4. Logic Based AI

IQ, reasoning...

e.g. Knowledge based system that can answer queries by applying logical rules. If you say 'All students are smart'

'John is a student' then it'll infer 'John is smart'

2. Natural Language Processing

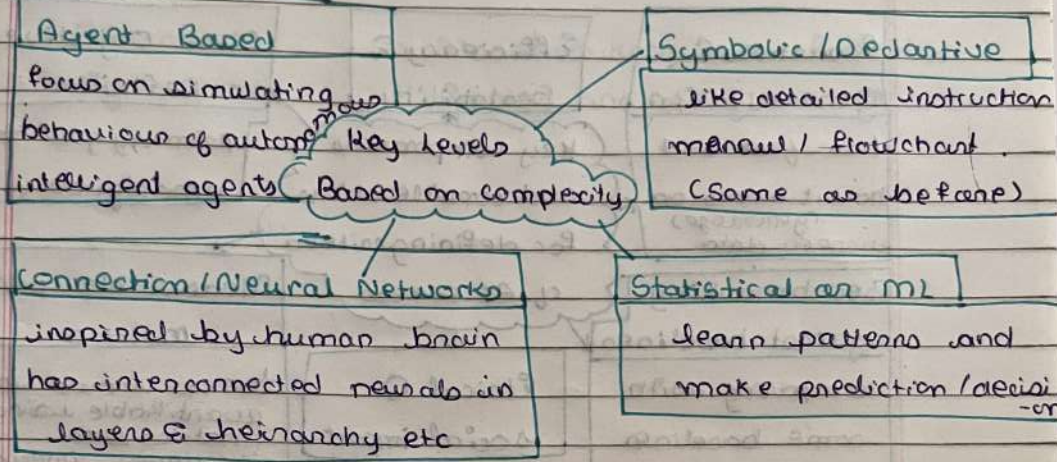
- enables computers to understand, interpret and generate human language
- basically allows AI to understand what human says and responds in human way...

3. Computer Vision

- allows computers to see, process and understand digital images and videos from the real world.
- relies on ML techniques to identify objects, recognize patterns and interpret scenes.
eg. facial recognition system
Automated quality inspection system using cam.

THE LEVEL OF MODEL

- describe how advanced or capable an AI system is, based on how much it can think, learn and act like humans
- the level helps us to categorize these models based on what they aim to capture and how they respond / act / operate.

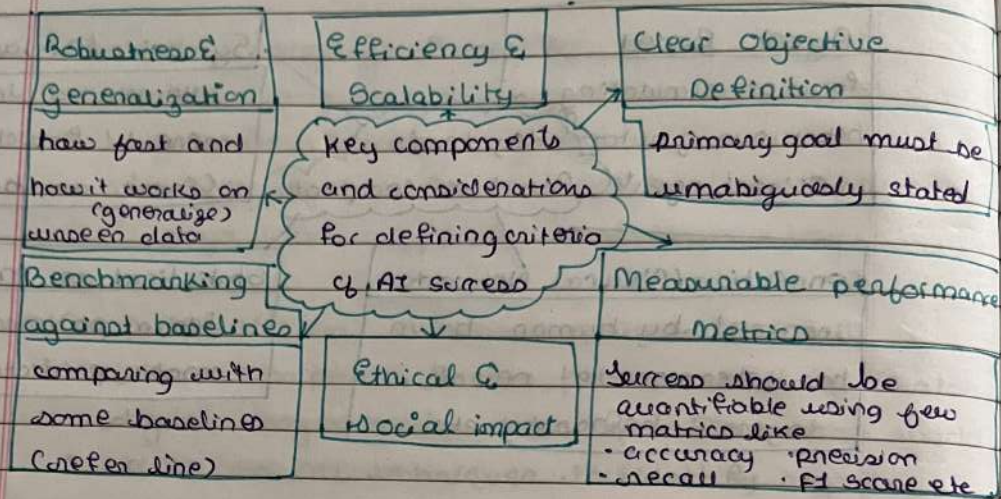


others

Artificial Narrow Intelligence	good at tricks, only a trained dog which one taught
Artificial General Intelligence	A human can learn anything knows everything
Artificial Super Intelligence	A god like being better than human

CRITERIA FOR SUCCESS

- when developing an AI, it is crucial to define what success looks like before, during and after its creation
- It helps evaluating AI, measuring performance etc



APPLICATION OF AI

1. Personal Assistant and Smart Devices
2. Healthcare
3. Financial and Banking
4. Natural Language Processing Applications
 - ↳ Machine Translation
 - ↳ Spam Filtering
 - ↳ Sentiment analysis
5. Computer Vision Applications
 - ↳ Face recognition
 - ↳ Quality control
 - ↳ Object detection
6. Transportation and Autonomous vehicles
7. Ecommerce and retail

8. Education

9. Agriculture

10. Manufacturing and Robotics

11. Gaming and Entertainment