

Notes on: State Space Search and Heuristic Technique_from_0

1.) Solving problems as state space search

Solving Problems as State Space Search

In Artificial Intelligence, many complex problems can be modeled and solved efficiently using a technique called **state space search**. This method transforms a problem into a representation that an AI agent can explore systematically to find a solution. The core idea is to see the problem as navigating through a landscape of possible situations.

The Core Idea: Representing a Problem as a State Space

Imagine a problem as a collection of possible configurations or **states**. You start at one state (the initial state) and want to reach another (the goal state) by performing valid actions. The entire collection of these possible states and the connections between them forms the **state space**.

Key Components of State Space Search:

1. State:

- Definition: A **state** is a complete, unambiguous description of the problem at a specific moment. It contains all the information needed to decide what actions are possible and what the current situation is.
- Example: In a simple 8-puzzle game, a state would be the exact arrangement of the tiles on the 3x3 board. For a robot navigating a room, a state could be its current (x, y) coordinates and orientation.

2. State Space:

- Definition: The **state space** is the set of all reachable states from the initial state through any sequence of valid actions. It represents every possible configuration or situation the problem can be in.
- Example: For the 8-puzzle, the state space includes all possible unique tile arrangements (there are $9!/2 = 181,440$ such states). For a map, it includes all cities or junctions.

3. Operators (or Actions/Transitions):

- Definition: **Operators** are actions or moves that transform one state into another. Each operator defines what changes occur and often has preconditions (conditions that must be met to apply the operator) and effects (the result of applying it).
- Example: In the 8-puzzle, sliding an adjacent tile into the empty space is an operator. For route planning, taking a road from city A to city B is an operator.

4. Initial State:

- Definition: This is the starting configuration of the problem, the state you begin in before any actions have been taken.
- Example: The scrambled tile arrangement of an 8-puzzle before any moves. Your current location on a map when you start planning a trip.

5. Goal State (or Goal Test):

- Definition: This is the desired final state(s) or a condition that, when met, signifies that the problem has been solved. Some problems have multiple goal states, requiring a **goal test** function to check if any current state is a goal.
- Example: The sorted tile arrangement (e.g., 1-2-3-4-5-6-7-8-blank) of an 8-puzzle. Arriving at a specific destination city in a navigation problem.

The Search Process: Finding a Solution Path

Solving a problem as state space search means finding a sequence of operators that leads from the initial state to a goal state. This sequence is known as the **solution path**.

- The process involves systematically exploring the state space. It starts from the initial state, applies operators to generate new states, and continues this exploration until a goal state is discovered.
- The term **search** refers to the systematic exploration algorithms used to navigate this vast space of possibilities.

Why This Representation is Powerful:

- **Standardization:** It offers a universal and consistent framework for representing a wide variety of problems, making them amenable to automated solutions.
- **Foundation for Algorithms:** This representation forms the basis upon which various search algorithms (like Depth-First Search, Breadth-First Search, Best-First Search) are built to systematically explore the problem space and find solutions.
- **Clarity:** It helps in breaking down complex problems into discrete, manageable components (states and transitions), making them easier to analyze and understand.

Real-World Example - Finding a Route on a Map:

- **State:** Your current location (e.g., a city, a specific intersection).
- **State Space:** All possible cities, towns, and junctions on the map.
- **Operators:** Traveling along a road from one location to an adjacent location.
- **Initial State:** Your starting city.
- **Goal State:** Your destination city.
- **Solution:** The sequence of roads (operators) and cities (states) that takes you from the start to the destination.

This state space model is fundamental in AI because it allows agents to methodically explore potential actions and their outcomes to achieve objectives, laying the groundwork for more advanced problem-solving techniques.

Summary of Key Points:

- **State space search** is a core AI problem-solving method.
- Problems are modeled using States, Operators, an Initial State, and a Goal State.
- **State:** A snapshot of the problem's current situation.
- **State Space:** The collection of all possible states.
- **Operators:** Actions that change one state into another.
- **Initial State:** The starting configuration.
- **Goal State:** The desired target configuration or condition.
- The solution is a specific path (sequence of operators) through this state space.
- This framework enables AI agents to systematically search for solutions.

2.) Production system

In Artificial Intelligence, after understanding how problems can be modeled as state space searches, we need a mechanism to implement the process of moving through these states. A **Production System** provides a general and powerful framework for this.

1- What is a Production System?

A production system is a type of AI architecture or framework commonly used for implementing rule-based systems, expert systems, and for solving problems that can be represented as state-space search problems. It consists of a set of components that work together to reason and make decisions.

2- Core Components of a Production System

A production system typically has three main components:

2.1- Global Database (or Working Memory)

- This is the central data structure that holds all the current information or facts about the problem at hand.
- It represents the current state of the world or the problem.
- Information in the global database can be modified by the production rules.
- Example: In a robot navigation problem, the database might contain **robot_at(x,y)**, **obstacle_at(x+1,y)**, **goal_at(gx,gy)**. In a game like Chess, it would be the current board configuration.

2.2- Production Rules (or Knowledge Base)

- These are the heart of the system, representing the knowledge.
 - Each rule is typically in an **IF-THEN** or **CONDITION-ACTION** format.
 - IF part (LHS - Left Hand Side): A set of conditions that must be true for the rule to be applicable.
 - THEN part (RHS - Right Hand Side): A set of actions to be performed if the conditions are met.
- These actions modify the global database, effectively changing the state.
- Example Rules:
 - IF **robot_at(X,Y)** AND **no_obstacle_ahead** THEN **move_forward(X,Y)**
 - IF **robot_at(X,Y)** AND **obstacle_ahead** THEN **turn_right**
 - IF **current_patient_has_fever** AND **current_patient_has_cough** THEN **diagnose_possible_flu**

2.3- Interpreter (or Control Strategy / Inference Engine)

- This is the brain that decides which rule to apply and when.
 - It manages the execution of the production rules.
 - Its primary task is the **recognize-act cycle**:
- 1- Recognize: It scans the global database to find all production rules whose IF-part conditions are satisfied by the current state. These are the **active** or **conflict set** rules.
 - 2- Conflict Resolution: If multiple rules are active, the interpreter uses a strategy to select only one rule to execute. Strategies might include:
 - First rule encountered.
 - Rule with more specific conditions.
 - Rule that refers to the most recently added facts.
 - Rule with higher priority.
 - 3- Act: It executes the THEN-part actions of the selected rule. This modifies the global database, leading to a new state.
 - This cycle repeats until a goal state is reached, no rules are applicable, or a predefined termination condition is met.

3- Production Systems in State Space Search Context

- In state space search, the **global database** represents the current state of the problem.
- The **production rules** define the valid **operators** or **actions** that can transition the system from one state to another.
- The **interpreter** is responsible for applying these operators (rules) to explore the state space, moving from the initial state towards a goal state.
- Each application of a rule effectively generates a new state from the current state.
- The choice of which rule to apply next is crucial and often guided by search algorithms (like Depth First Search, Breadth-First Search, Best First Search) and heuristic functions, which are topics you will cover later.

4- Real-world Analogy

Imagine a cooking recipe book.

- Global Database: Your kitchen counter with ingredients and current status (e.g., **onions chopped**, **stove hot**). This is your current **state**.
- Production Rules: Each step in the recipe is a rule. **IF 'onions chopped' AND 'pan hot' THEN 'add onions to pan and sauté for 5 minutes'**. This is an **operator** that changes your kitchen state.
- Interpreter: You, the cook, reading the recipe, checking if conditions are met, deciding which step to do next (e.g., if multiple steps can be done simultaneously), and performing the action. You are guiding the process from raw ingredients (initial state) to a finished meal (goal state).

Summary of Key Points:

- A Production System is an AI architecture for rule-based problem-solving and state-space search.
- It comprises a Global Database (current state/facts), Production Rules (IF-THEN knowledge/operators), and an Interpreter (control strategy).
- The Interpreter operates in a **recognize-act cycle** to select and apply rules, modifying the global database and moving between states.
- It provides a structured way to define states, valid transitions (operators), and a control mechanism to explore the problem's state space.

3.) Problem characteristics

Problem characteristics are essential features of a problem that help us understand its nature and, more importantly, guide us in choosing the most suitable search strategy or algorithm to solve it within a state space. When a problem is defined as a state space, its characteristics define the structure of that space and the challenges of navigating it.

Why Problem Characteristics Matter:

- They determine which search algorithms (like those you'll learn later - Depth-First Search, Breadth-First Search, Hill Climbing, Best-First Search) are effective.
- They help assess the complexity and resource requirements (time, memory) for solving a problem.
- They inform the design of problem-solving agents.

Here are the key characteristics of a problem:

1- Is the problem decomposable?

- Explanation: Can the problem be broken down into smaller, independent subproblems, each of which can be solved separately? If so, the solutions to the subproblems can then be combined to form the solution to the original problem.
- Example: Sorting a large list of numbers can be decomposed into sorting smaller sub-lists and then merging them (e.g., Merge Sort). Proving a complex mathematical theorem can be broken into proving several lemmas.
- Impact: Decomposable problems can often be solved more efficiently by tackling parts individually, sometimes even in parallel.

2- Can solution steps be ignored, undone, or must they be irreversible?

This characteristic describes the permanence of actions within the problem's state space.

- Ignorable:
 - Explanation: If a step taken turns out to be wrong, it can simply be ignored or forgotten, and another path can be tried without any significant cost or consequence.
 - Example: Proving a theorem – if one line of reasoning doesn't lead to a solution, you can simply abandon it and try another approach without having to **undo** anything.
 - Impact: Simplifies search as no complex backtracking mechanisms are needed.
- Recoverable:
 - Explanation: If a step proves incorrect, it's possible to backtrack and undo the move to return to a previous state, or to take compensatory actions. There's a cost, but mistakes aren't permanent.
 - Example: Playing chess (in a casual game, you might retract a move) or solving a Rubik's Cube (you can always reverse a bad turn).
 - Impact: Requires algorithms that can manage state history and enable backtracking.
- Irreversible:
 - Explanation: Once a step is taken, it cannot be undone or reversed. Mistakes have permanent consequences.
 - Example: A robot arm picking up a fragile object – dropping it cannot be undone. Performing surgery or constructing a building.
 - Impact: Demands very careful planning and forward-looking search algorithms that minimize errors,

as backtracking is impossible or too costly.

3- Is the universe predictable? (Certainty vs. Uncertainty)

- Explanation: Do we know the exact outcome of every action we take?
- Certainty: The outcomes of all actions are fully known and deterministic.
- Example: Playing a game like Sudoku – when you place a number, its effect on the board is certain.
- Uncertainty: The outcomes of actions are not fully known or might be probabilistic. External factors can influence the result.
- Example: A robot navigating in a real-world environment – its sensors might be noisy, or its movements might not be perfectly precise.
- Impact: Certain problems can use deterministic planning; uncertain problems require more complex strategies, often involving sensing and re-planning.

4- Is a good solution absolute or relative? (Satisficing vs. Optimality)

- Explanation: Are we looking for **any** solution that meets criteria, or the **best possible** solution according to some measure?
- Satisficing: Finding a solution that is **good enough** or meets a minimum set of requirements, even if it's not the absolute best.
- Example: Finding a route to a restaurant that takes less than 20 minutes, even if a 15-minute route exists.
- Optimality: Finding the best possible solution according to a defined cost function (e.g., shortest path, lowest cost, fastest time).
- Example: Finding the **absolute shortest** route to a restaurant.
- Impact: Optimal solutions often require more extensive (and thus slower) search, exploring more of the state space. Satisficing can be much faster.

5- Is the solution a state or a path?

- Explanation: Is the goal simply reaching a particular state, or is the sequence of actions taken to reach that state also part of the solution?
- State Solution: The final state itself is the answer. The steps to get there might not be explicitly required.
- Example: Proving a theorem – the theorem itself is the solution; the particular steps of the proof might not be the primary output (though useful).
- Path Solution: The sequence of actions (the path) from the initial state to the goal state is the required output.
- Example: Robot navigation – the robot needs to know the **sequence of movements** to get from point A to point B.
- Impact: Algorithms for path solutions need to store and reconstruct the path, which can add memory requirements.

6- What is the role of knowledge? (With or Without Knowledge)

- Explanation: Does the problem provide additional information beyond its basic definition that can guide the search?
- Without Knowledge (Uninformed Search / Blind Search): The problem description itself provides no special information to help prioritize which states to explore next. Search proceeds systematically (e.g., trying all neighbors).
- Example: Solving a maze without seeing the layout, just trying paths.
- Impact: Generally less efficient for large state spaces. (Relevant to DFS, BFS).
- With Knowledge (Informed Search / Heuristic Search): The problem provides additional knowledge (often in the form of a heuristic function) that estimates how close a state is to the goal, guiding the search more directly.
- Example: Solving a maze by always trying paths that lead **generally towards** the exit.
- Impact: Can significantly speed up search by focusing on promising paths. (Relevant to Heuristic function, Hill Climbing, Best First Search).

7- Does the problem require interaction? (Online vs. Offline)

- Explanation: Is the entire solution planned before any action is taken, or is planning interleaved with acting and sensing the real world?
- Offline: All calculations and planning are completed before any execution begins. The agent knows everything upfront.

- Example: Solving a Sudoku puzzle on paper. All decisions are made before writing anything down.
- Online: The agent plans some actions, executes them, observes the results from the environment, and then plans the next set of actions. This is common in real-world, uncertain environments.
- Example: A self-driving car – it plans a short segment, drives, senses traffic changes, and re-plans.
- Impact: Online problems are more complex, needing to handle real-time constraints, sensing, and uncertainty.

Summary:

Problem characteristics help categorize problems and select appropriate search strategies. They tell us if a problem can be broken down, how forgiving it is to mistakes, if outcomes are certain, whether we need the best solution or just a good one, what the solution looks like, if extra guidance is available, and if real-time interaction is needed. Understanding these features is fundamental to effectively applying AI search techniques.

4.) Depth First Search

Depth First Search (DFS)

Recall that in Artificial Intelligence, solving problems often involves navigating a 'state space' – a conceptual graph where each node represents a possible state of the problem and edges represent actions that transition from one state to another. Finding a solution means finding a path from an initial state to a goal state within this space. Depth First Search (DFS) is one fundamental strategy to explore such a state space. It is an uninformed search algorithm, meaning it doesn't use any domain-specific knowledge or heuristics to guide its search, only the structure of the state space itself.

1- What is Depth First Search (DFS)?

Depth First Search is a search strategy that explores as far as possible along each branch before backtracking. Imagine you are in a maze; DFS is like picking a path, following it until you hit a dead end, and only then going back to the last junction to try another path. It prioritizes going deep into a single path rather than exploring all immediate options.

2- How DFS Works (Algorithm/Process)

The core idea of DFS involves a systematic exploration:

- Start at the initial state (often called the 'root' node of the search tree).
- Choose one of its unvisited successor states (children).
- Mark the current state as visited.
- Now, make this chosen successor state the new current state and repeat the process: explore one of its unvisited successors. This continues, pushing deeper into the state space.
- If you reach a state that has no unvisited successors (a dead end) or if the state has already been visited on the current path (to prevent cycles), then 'backtrack'. Backtracking means returning to the most recent state that still has unvisited successors and exploring one of those.
- This process continues until the goal state is found or until all reachable states have been visited.
- A stack data structure is implicitly or explicitly used to keep track of the path taken and the nodes to visit next, allowing for efficient backtracking.

3- Characteristics of DFS

- **Completeness:** DFS is not guaranteed to find a solution if one exists in an infinite state space, as it might get stuck going infinitely deep down a fruitless path. However, in finite state spaces and with proper cycle detection (avoiding revisiting nodes on the current path), it is complete.
- **Optimality:** DFS is generally not optimal. The first solution it finds might not be the shortest or the least-cost path. It could find a very deep, long path to the goal before discovering a much shorter one.
- **Time Complexity:** In the worst case, DFS might have to visit all nodes in the state space. If 'b' is the branching factor (average number of successors per state) and 'm' is the maximum depth of the search tree, the time complexity is $O(b^m)$.
- **Space Complexity:** DFS has a significant advantage in space complexity. It only needs to store the current path from the root to the current node, plus the unvisited siblings of the nodes on that path. This

makes its space complexity $O(b \cdot m)$, or even $O(m)$ if only the current path is stored and successors are generated on the fly. This is generally much less than BFS for deep trees.

4- Advantages and Disadvantages of DFS

- Advantages:
 - Memory Efficiency: For very deep state spaces, DFS uses significantly less memory than Breadth-First Search (BFS) because it doesn't need to store all nodes at a given level.
 - Can Find Solutions Quickly: If a solution exists deep down a particular path, DFS can find it much faster than BFS, which explores level by level.
 - Useful for Generating Specific Paths: When you need to find *any* path to a goal, and not necessarily the shortest, DFS can be effective.
- Disadvantages:
 - Can Get Trapped: Without proper mechanisms to detect visited states, DFS can get caught in infinite loops in cyclic graphs or explore infinitely deep paths in infinite state spaces without ever finding a solution.
 - Not Optimal: It does not guarantee finding the shortest or optimal path, as it might explore a long path to a solution first.
 - Can Miss Shallow Solutions: If the goal is at a shallow depth but down a path that is not explored first, DFS might spend a lot of time exploring a deep, fruitless branch before eventually backtracking to find the shallow solution.

5- Real-World Examples and Applications

- Maze Solving: DFS is intuitively good for solving mazes. You follow one path until a dead end, then backtrack and try another.
- Finding Connected Components: In graph theory, DFS can be used to determine all nodes reachable from a starting node, thus finding connected components.
- Topological Sorting: Ordering tasks or events where some must precede others (e.g., in project scheduling).
- Web Crawlers: Some specialized web crawlers might use a DFS-like approach to explore deeply into a specific section of a website.
- Pathfinding in Games: While often optimized with heuristics, the basic concept of exploring a path fully before turning back is present in some AI pathfinding for game agents.

6- Summary of Key Points

- DFS explores a state space by going as deep as possible along a single path.
- It backtracks when it hits a dead end or a previously visited node.
- It is an uninformed search strategy, meaning it doesn't use heuristics.
- It is memory-efficient, especially for deep solutions, due to its stack-based approach.
- It is generally not complete in infinite spaces and not optimal (doesn't guarantee the shortest path).
- Useful in applications like maze solving, graph traversal, and topological sorting.

5.) Breadth-First Search

Breadth-First Search (BFS)

Breadth-First Search is a fundamental uninformed search algorithm used in Artificial Intelligence to explore a state space. It systematically explores all reachable states level by level, ensuring that it fully examines all states at a given depth before moving to the next deeper level. This contrasts with Depth-First Search (DFS) which explores as deeply as possible along one path before backtracking.

1. Core Idea and Mechanism

- Level-by-level exploration: BFS expands the search by examining all immediate successor states (level 1) of the initial state, then all their successors (level 2), and so on.
- Queue data structure: BFS uses a Queue (First-In, First-Out, or FIFO) to manage the order in which nodes are visited. New successor nodes are added to the back of the queue, and nodes for expansion are taken from the front.

- Visited set: It maintains a set of visited nodes to prevent cycles and avoid re-processing states, ensuring efficiency.

2. How Breadth-First Search Works (Algorithm Intuition)

Imagine you are exploring a maze. Instead of picking one path and following it until you hit a dead end or find the exit (like DFS), BFS explores all possible paths one step at a time. It first checks all adjacent cells, then all cells reachable in two steps, then three steps, and so on, radiating outwards from the start.

3. Steps of the BFS Algorithm

- Initialize a Queue and add the initial state (root node).
- Create a 'Visited' set and add the initial state to it.
- While the Queue is not empty:
 - Dequeue a node (let's call it `current_node`).
 - If `current_node` is the goal state, the search is successful; return the path.
 - Generate all valid successor states (children) of `current_node`.
 - For each successor:
 - If the successor has not been visited:
 - Mark it as visited.
 - Enqueue it.

4. Characteristics and Properties

- Completeness: BFS is complete, meaning if a solution (goal state) exists in the state space, BFS is guaranteed to find it. This is because it systematically explores every reachable state.
- Optimality: For unweighted graphs or problems where all step costs are equal (e.g., each move costs 1), BFS is optimal. It guarantees finding the shortest path (in terms of number of steps) from the initial state to the goal state.
- Time Complexity: $O(b^d)$, where 'b' is the branching factor (average number of successors for a node) and 'd' is the depth of the shallowest goal state. It might need to explore all nodes up to depth 'd' in the worst case.
- Space Complexity: $O(b^d)$. This is the main disadvantage. BFS needs to store all nodes at the current level in the queue, which can grow exponentially with depth.

5. Real-World Examples

- Shortest Path in Networks: Finding the minimum number of hops (connections) between two computers in a network or two people in a social media graph.
- Web Crawlers: Search engine bots use BFS-like approaches to discover new web pages by following links from initial seed pages.
- Pathfinding in Games: Finding the shortest path for a character in a grid-based game where each move has equal cost.

6. Advantages of BFS

- Guaranteed to find a solution if one exists.
- Guaranteed to find the optimal solution (shortest path) in problems with uniform step costs.
- Good for shallow goal states or when the depth of the solution is unknown.

7. Disadvantages of BFS

- High memory consumption: Can quickly exhaust memory for large state spaces or deep solutions due to storing all nodes at the current level.
- Can be very slow: If the branching factor is large, it may explore a huge number of states before reaching a deep goal.
- Not suitable for problems with extremely deep solutions or limited memory resources.

8. Summary of Key Points

- BFS explores states level by level, always expanding the shallowest unexpanded node first.
- It uses a Queue (FIFO) to manage node expansion order.
- BFS is complete and optimal for unweighted graphs or uniform step costs.
- Its main drawback is its high space complexity, making it challenging for very large state spaces.

6.) Heuristic function

1. Introduction to Heuristic Functions

In Artificial Intelligence, many problems are solved by exploring a **state space** - a graph where nodes represent possible states and edges represent actions to move between states. Our goal is to find a path from an initial state to a target or **goal** state. Uninformed search methods like Depth-First Search (DFS) and Breadth-First Search (BFS) explore this space systematically but without any special knowledge about where the goal might be. For complex problems with very large state spaces, these methods can be inefficient, taking too long or requiring too much memory. This is where **heuristic functions** come in.

2. What is a Heuristic Function?

A heuristic function, often denoted as $h(n)$, is an estimation function that takes a problem state 'n' as input and returns a numerical value. This value is an estimate of the **cost** or **distance** from the current state 'n' to the nearest goal state.

- It doesn't guarantee the best path, but it provides a **guess** or **rule of thumb** to guide the search towards the goal.
- Think of it like a shortcut or an educated guess to find your way in a new city without a perfect map, but with some sense of direction.

3. Why do we need Heuristic Functions?

We need heuristic functions to improve the efficiency of search algorithms in large state spaces.

- **Guidance:** They help direct the search towards promising paths, avoiding fruitless exploration.
- **Efficiency:** They allow search algorithms to find solutions much faster than uninformed methods, especially for complex problems.
- **Feasibility:** For problems with extremely large or infinite state spaces, an uninformed search might never find a solution. Heuristics make such problems tractable.

4. Properties of a Good Heuristic

- **Informativeness:** A good heuristic provides an estimate that is as close as possible to the true cost. The more informed it is, the better it guides the search.
- **Computability:** It should be easy and fast to compute the heuristic value for any given state. If computing $h(n)$ takes too long, it defeats the purpose of speeding up the search.
- **Admissibility:** A heuristic is **admissible** if it never overestimates the actual cost to reach the goal from any state. That is, $h(n) \leq \text{true_cost}(n, \text{goal})$. Admissible heuristics are crucial for guaranteeing optimal solutions in algorithms like A^* .
- **Consistency (or Monotonicity):** A heuristic is **consistent** if for every node n and every successor n' of n , the estimated cost from n to the goal is no greater than the cost of moving from n to n' plus the estimated cost from n' to the goal. That is, $h(n) \leq \text{cost}(n, n') + h(n')$. Consistent heuristics imply admissibility and simplify pathfinding in some algorithms.

5. How Heuristic Functions Work (Examples)

Consider the 8-Puzzle problem (a 3x3 grid with 8 numbered tiles and one blank space, objective is to arrange tiles in numerical order).

- State: A particular arrangement of tiles.
- Goal: The target arrangement.

Two common heuristic functions for the 8-puzzle:

a) Number of Misplaced Tiles (h_1):

- This heuristic counts how many tiles are not in their correct goal positions.
- Example: If the goal is [1 2 3 / 4 5 6 / 7 8 _] and current state is [1 2 3 / 4 5 _ / 7 8 6], then tile '6' is misplaced. $h_1 = 1$.
- It's an admissible heuristic because each misplaced tile needs at least one move to get to its correct position, and it never overestimates the true moves needed.

b) Manhattan Distance (h_2):

- This heuristic calculates the sum of the horizontal and vertical distances (number of grid squares) each tile is away from its correct goal position. The blank tile is ignored.
- Example: For the tile '6' in the example above, if its goal position is (2,3) and its current position is (3,3), its Manhattan distance is $|3-2| + |3-3| = 1$.
- Manhattan distance is generally a more informed and powerful heuristic than the number of misplaced tiles for the 8-puzzle, and it is also admissible.

6. Real-World Applications

- Pathfinding in GPS and navigation systems: Estimating the remaining travel time or distance to a destination.
- Game AI: Guiding game characters to a target, estimating the best move in complex board games (e.g., chess, checkers).
- Logistics and supply chain: Finding efficient routes for delivery vehicles.
- Robotics: Planning movements and actions for robots to achieve tasks.

7. Summary

- A heuristic function $h(n)$ estimates the cost from a current state 'n' to the goal state.
- It guides search algorithms to be more efficient than uninformed methods.
- Good heuristics are informative, computationally efficient, and ideally admissible or consistent to guarantee optimal solutions.
- They are essential tools for solving complex AI problems in practical applications.
- Heuristic functions form the basis for **informed search** algorithms, which you will encounter in topics like Best-First Search.

7.) Hill climbing

Hill Climbing is a local search algorithm used in Artificial Intelligence (AI) to solve optimization problems. It falls under the category of State Space Search, but unlike exhaustive search methods like Depth First Search (DFS) or Breadth-First Search (BFS) that explore the entire state space broadly, Hill Climbing is a greedy approach. It uses a heuristic function to guide its search for a solution by continuously moving towards a better state.

Imagine you are climbing a hill in the fog, and you want to reach the highest point. You can only feel the slope immediately around you. Your strategy would be to always take a step in the direction that goes upwards the most. You keep doing this until you can't find any direction that goes further up. This is the core idea of Hill Climbing.

How Hill Climbing Works:

1- Start: Begin at an arbitrary or randomly chosen initial state within the state space. This state represents a possible solution or configuration.

2- Evaluate Current State: Use a heuristic function to calculate the **value** or **goodness** of the current state. For optimization problems, this value indicates how close the state is to the goal. For example, if minimizing a cost, a lower value is better. If maximizing profit, a higher value is better.

3- Generate Neighbors: Identify all possible **neighbor** states that can be reached from the current state by making a small change (e.g., one move, one alteration). These are the states immediately adjacent to the current one in the state space.

4- Evaluate Neighbors: Calculate the heuristic value for each of these neighbor states.

5- Move:

- If any neighbor state has a better heuristic value than the current state, move to the best among those better neighbors.
- If no neighbor state has a better heuristic value than the current state, the algorithm stops. It has reached a **peak** (a local maximum if maximizing, or a local minimum if minimizing).

6- Repeat: Continue steps 2-5 from the new current state until a peak is reached.

Types of Hill Climbing:

- Simple Hill Climbing: It moves to the first neighbor state that is better than the current state. It doesn't check all neighbors. It's faster but less thorough.
- Steepest Ascent Hill Climbing: It examines all neighbor states and moves to the neighbor state that offers the greatest improvement in the heuristic value. It's slower per step but often makes more progress.
- Stochastic Hill Climbing: Instead of choosing the best neighbor, it chooses a random neighbor from the **better** neighbors (neighbors that improve the current state). The probability of choosing a neighbor might depend on the degree of improvement.

Challenges and Limitations (Why it's a Local Search):

Hill Climbing is a greedy algorithm, meaning it only considers the immediate best step. This often leads to several problems:

1- Local Maxima (or Minima): This is the most significant limitation. The algorithm might reach a state that is better than all its immediate neighbors but is not the globally best state (the true highest point). Imagine reaching a small hill's summit, thinking it's the mountain's highest point, while a much taller peak exists further away. It gets stuck here because all steps lead downwards.

2- Ridges: A ridge is a sequence of local maxima that are hard to navigate. The algorithm might struggle to make progress along a ridge because each move horizontally along the ridge might not show an immediate improvement in height, or it might require multiple steps that aren't individually **upward**.

3- Plateaus: A plateau is a flat area in the search space where the heuristic function yields the same value for all neighboring states. The algorithm cannot determine which direction to move to find a better state and may wander aimlessly or get stuck.

Advantages of Hill Climbing:

- Simplicity: Easy to understand and implement.
- Memory Efficiency: It only needs to store the current state and its neighbors, making it very memory-efficient compared to algorithms that explore large parts of the state space.
- Good for Specific Problems: Effective when the search space is large, and finding a **good enough** solution quickly is more important than finding the absolute best solution.

Real-World Examples and Applications:

- Automated Design: Optimizing parameters in engineering design (e.g., circuit design, antenna placement).
- Scheduling Problems: Finding efficient schedules for tasks or resources.
- Network Optimization: Adjusting network configurations to improve performance.
- Machine Learning: Gradient descent, a fundamental optimization algorithm used to train neural networks, is conceptually similar to Hill Climbing, where the **hill** is the loss function, and we are trying to find its minimum.

Summary of Key Points:

- Hill Climbing is a local, greedy search algorithm.
- It uses a heuristic function to move from the current state to a better neighbor state.
- It continues until no better neighbor can be found, reaching a local optimum.
- Key limitations include getting stuck in local maxima, ridges, and plateaus.
- It is simple, memory-efficient, and useful for finding satisfactory solutions quickly in large state spaces.
- Its limitations highlight the need for more sophisticated heuristic search techniques that try to overcome these pitfalls.

8.) Best First Search

Best First Search

Best First Search (BFS) is a type of informed search algorithm used in Artificial Intelligence for navigating state spaces to find a path from a start state to a goal state. It belongs to the family of algorithms that use heuristic information to guide their search, making them generally more efficient than uninformed search methods like Depth-First Search (DFS) or Breadth-First Search (BFS – a different one, Breadth-First Search).

1- What is Best First Search?

- It's a search strategy that always expands the node which appears **best** according to some evaluation function, typically a heuristic function.
- Unlike uninformed searches that explore systematically (like level by level or path by path), Best First Search prioritizes states that seem closer to the goal.

2- How it Works

- Best First Search combines the advantages of Depth-First Search (following a path deeply) and Breadth-First Search (exploring neighbors) by using a heuristic to make informed decisions.
- It maintains a list of unexpanded nodes (often called the 'Open List' or 'Frontier') and a list of already expanded nodes (the 'Closed List' or 'Visited List').
- The 'Open List' is implemented as a priority queue, where nodes are ordered based on their heuristic evaluation.

3- Key Components

- Heuristic Function ($h(n)$): Recall that a heuristic function estimates the cost or distance from the current node 'n' to the goal state. A good heuristic helps guide the search effectively.
- Priority Queue (Open List): This data structure stores all generated but not yet expanded nodes. Nodes are retrieved from it based on their heuristic value, always picking the **best** (usually the one with the lowest estimated cost to the goal).
- Closed List: Keeps track of nodes that have already been expanded to prevent redundant processing and cycles.

4- Greedy Best First Search (GBFS)

- Greedy Best First Search is the most common implementation of Best First Search.
- Its evaluation function for a node 'n' is simply the heuristic function $h(n)$.
- It always expands the node that has the lowest $h(n)$ value, meaning the node that is estimated to be closest to the goal.

• Example Analogy: Imagine you're driving to a new city without a precise GPS, but you have a general sense of where the city is. At every intersection, you choose the road that seems to point most directly towards the city, based on landmarks or a rough mental map. You don't consider how many turns are on that road or how long it actually is, just its apparent direction. This is like Greedy Best First Search – it always takes the path that **looks** best at that moment based on the heuristic.

5- Algorithm Steps for Greedy Best First Search

1. Initialize the Open List with the start node.
2. Initialize the Closed List as empty.
3. While the Open List is not empty:
 - a. Remove the node 'n' from the Open List with the lowest $h(n)$ value.
 - b. If 'n' is the goal state, then a solution has been found. Reconstruct and return the path.
 - c. Add 'n' to the Closed List.
 - d. Generate all successor nodes of 'n'. For each successor 's':
 - i. If 's' is not in the Open List and not in the Closed List, calculate $h(s)$ and add 's' to the Open List.
 - ii. If 's' is already in the Open List or Closed List, it's typically ignored in the simplest GBFS, or its path might be re-evaluated (though this is more common in A*).

6- Advantages

- Efficiency: Often finds a solution much faster than uninformed search methods, especially in large state spaces, because it's directed towards the goal.
- Space Complexity: Can be more space-efficient than Breadth-First Search if the heuristic is good.

7- Disadvantages

- Not Optimal: Greedy Best First Search is not guaranteed to find the shortest or optimal path. It might make a locally optimal choice that leads to a longer path overall.
- Not Complete: It can get stuck in infinite loops if it repeatedly expands nodes that lead back to previously visited states, or if the heuristic guides it away from the goal in a misleading way. Using a Closed List helps mitigate this, but it can still get trapped if the optimal path is **hidden** behind poor heuristic values.
- Susceptible to misleading heuristics: If the heuristic function is inaccurate, the search can be very inefficient or fail to find a solution.

8- Real-World Applications

- Pathfinding in games: Often used for quick, good-enough pathfinding for non-player characters where finding the absolute shortest path isn't critical but speed is.
- Robotics: For navigation, to quickly estimate a path to a target without complex computations.
- Logistics and planning: For finding reasonable solutions quickly in complex routing problems.

Summary of Key Points:

- Best First Search uses a heuristic function to guide its search, prioritizing nodes estimated to be closest to the goal.
- It employs a priority queue (Open List) to store and retrieve the **best** nodes to expand.
- Greedy Best First Search is a common variant where the evaluation function is solely the heuristic value $h(n)$.
- It is generally more efficient than uninformed searches but is neither optimal (doesn't guarantee the shortest path) nor complete (might not find a solution in all cases).
- Despite its drawbacks, it's valuable for quickly finding good solutions in complex state spaces.

