

FINAL QUESTION PAPER

1. 42. TensorFlow and PyTorch are primarily used for developing and deploying: (a) Relational databases (b) Traditional statistical models (c) Deep learning and neural networks (d) Web servers
2. 31. Which of the following is a common application of Machine Learning in the healthcare sector? (a) Predicting patient wait times (b) Diagnosing diseases from medical images (c) Managing hospital administrative tasks (d) Scheduling doctor appointments
3. 46. Which programming language is predominantly used for developing machine learning applications due to its extensive libraries and frameworks? a) Java b) C++ c) Python d) JavaScript
4. 47. Which Python library is primarily used for numerical operations and array manipulation, forming the base for many other machine learning libraries? a) Matplotlib b) NumPy c) Pandas d) SciPy
5. 27. The concept of 'feedback' is most central to which type of machine learning, often guiding an agent's learning process through rewards or penalties? a) Supervised Learning b) Unsupervised Learning c) Reinforcement Learning d) Batch Learning
6. 5. Which programming language is widely used for Machine Learning due to its extensive libraries like TensorFlow and scikit-learn? (a) Java (b) C++ (c) Python (d) R
7. 22. Which of the following is a key requirement for Supervised Learning algorithms to function effectively? a) The ability to explore an environment and receive feedback. b) A large amount of unlabeled data to discover hidden patterns. c) Clearly defined input features and corresponding target outputs (labeled data). d) Mechanisms to reduce the dimensionality of complex datasets.
8. 16. Which type of machine learning involves training a model on a dataset that includes both input features and their corresponding correct output labels? a) Unsupervised Learning b) Supervised Learning c) Reinforcement Learning d) Semi-supervised Learning
9. 6. A machine learning model that learns by interacting with an environment and receiving rewards or penalties is operating under which paradigm? (a) Supervised Learning (b) Unsupervised Learning (c) Reinforcement Learning (d) Clustering
10. 4. Identifying distinct customer segments from a large dataset without prior labels is a typical application of which Machine Learning paradigm? (a) Supervised Learning (b) Reinforcement Learning (c) Anomaly Detection (d) Unsupervised Learning
11. 19. A machine learning model that predicts a student's final exam score based on their attendance and homework submissions is performing which type of supervised learning task? a) Classification b) Clustering c) Regression d) Dimensionality Reduction
12. 36. Which type of machine learning involves training a model on a dataset that has already been labeled with the correct output? (a) Unsupervised Learning (b) Reinforcement Learning (c) Supervised Learning (d) Semi-supervised Learning
13. 55. Which of the following machine learning applications is typically categorized under Unsupervised Learning? a) Predicting house prices based on features like size and location. b) Classifying emails as spam or

not spam. c) Grouping customers into distinct segments based on their purchasing behavior. d) Recognizing faces in images with labeled identity data.

14. 32. In retail, Machine Learning is widely used for: (a) Tracking employee attendance (b) Optimizing product placement (c) Generating personalized product recommendations (d) Managing supply chain logistics

15. 29. An overview of human learning often highlights our ability to learn from examples (like Supervised Learning), but also to discover structure in the world without explicit instruction (like Unsupervised Learning). Which of the following best describes how machines mimic human learning by trial and error in dynamic environments? a) Supervised Classification b) Unsupervised Clustering c) Reinforcement Learning d) Feature Engineering

16. 49. Which of the following is a cloud-based platform that offers services and tools specifically designed for developing, deploying, and managing machine learning models? a) Microsoft Word b) Google Cloud AI Platform c) Adobe Photoshop d) Salesforce CRM

17. 39. Spam email detection, where emails are classified as "spam" or "not spam" based on historical labeled data, is a classic example of: (a) Unsupervised learning (b) Supervised learning (c) Reinforcement learning (d) Association rule mining

18. 48. Keras is a high-level API used for building and training deep learning models. It can run on top of which of the following backends? a) Scikit-learn b) Apache Spark c) TensorFlow d) PyTorch

19. 58. Identifying whether an incoming email is "spam" or "not spam" is an example of which type of supervised learning problem? a) Regression b) Clustering c) Classification d) Dimensionality Reduction

20. 7. What is the fundamental difference in the input data requirement between Supervised and Unsupervised Learning? (a) Supervised learning requires numeric data, while unsupervised learning requires categorical data. (b) Supervised learning requires labeled data, while unsupervised learning works with unlabeled data. (c) Supervised learning requires real-time data, while unsupervised learning uses historical data. (d) Supervised learning focuses on input features, while unsupervised learning focuses on output predictions.

21. 41. Which of the following is a popular Python library for numerical operations, commonly used as a foundation for machine learning algorithms? (a) Pandas (b) Matplotlib (c) NumPy (d) Scikit-learn

22. 38. A self-driving car learning to navigate by trial and error, receiving rewards for successful maneuvers and penalties for errors, best exemplifies: (a) Supervised Learning (b) Unsupervised Learning (c) Reinforcement Learning (d) Online Learning

23. 21. In Reinforcement Learning, what does an 'agent' learn through interaction with its environment? a) To label unlabeled data accurately. b) To predict a specific output given input features. c) An optimal policy of actions that maximizes a cumulative reward. d) To reduce the number of features in a dataset.

24. 52. Which interactive computing environment allows for the creation and sharing of documents that contain live code, equations, visualizations, and narrative text, making it popular for data exploration and machine learning experimentation? a) Integrated Development Environment (IDE) b) Command Line Interface (CLI) c) Jupyter Notebook d) Text Editor

25. 12. Provide two real-world applications where Machine Learning significantly improves decision-making or automation.

26. 37. Clustering customer segments based on their purchasing behavior, without prior labels, is an example of: (a) Supervised Learning (b) Unsupervised Learning (c) Reinforcement Learning (d) Transfer Learning

27. 45. Predictive maintenance in industrial settings, where machine failures are predicted before they occur using sensor data, typically falls under: (a) Unsupervised anomaly detection (b) Supervised time-series forecasting/classification (c) Reinforcement learning for control (d) Natural Language Processing

28. 40. Dimensionality reduction techniques like PCA are typically used in which type of machine learning? (a) Supervised Learning (b) Unsupervised Learning (c) Reinforcement Learning (d) Semi-supervised Learning

29. 3. Which type of Machine Learning is most appropriate for predicting house prices based on historical data including features like size, number of rooms, and location? (a) Unsupervised Learning (b) Reinforcement Learning (c) Supervised Learning (d) Semi-supervised Learning

30. 8. Define Machine Learning and explain its primary objective.

31. 26. What is the fundamental difference between Supervised and Unsupervised Machine Learning in terms of their learning objective? a) Supervised Learning aims to learn optimal actions; Unsupervised Learning aims to predict future events. b) Supervised Learning aims to predict or classify based on labeled data; Unsupervised Learning aims to find patterns in unlabeled data. c) Supervised Learning aims to interact with an environment; Unsupervised Learning aims to reduce data dimensions. d) Supervised Learning aims to discover hidden structures; Unsupervised Learning aims to assign data points to predefined categories.

32. 57. Which Python library is widely used for data manipulation and analysis, offering data structures like DataFrames for handling tabular data? a) TensorFlow b) PyTorch c) Pandas d) Scikit-image

33. 10. Categorize the following tasks into Supervised Learning, Unsupervised Learning, or Reinforcement Learning: (a) Predicting whether an email is spam or not spam. (b) Training a robot to navigate a maze. (c) Grouping news articles by topic.

34. 51. What specialized hardware component is crucial for accelerating the training of deep learning models due to its parallel processing capabilities? a) Central Processing Unit (CPU) b) Hard Disk Drive (HDD) c) Graphics Processing Unit (GPU) d) Random Access Memory (RAM)

35. 2. In the context of human learning, what is a key characteristic that enables adaptation and improvement over time? (a) Fixed set of rules for every situation. (b) Ability to generalize from past experiences. (c) Reliance solely on pre-programmed instructions. (d) Absence of error correction mechanisms.

36. 20. Discovering customer segments in a large dataset of purchase histories, without knowing the segments beforehand, is a typical application of: a) Supervised Classification b) Supervised Regression c) Reinforcement Learning d) Unsupervised Clustering

37. 56. In Reinforcement Learning, what is the primary role of the "reward signal"? a) To provide labeled examples for training the agent. b) To guide the agent on which actions are desirable or undesirable. c) To randomly select the next action for the agent. d) To store the agent's memory of past experiences.

38. 15. Name one common open-source library or framework used for implementing Machine Learning models and state its general purpose.

39. 35. Which aspect is generally more robust in human learning compared to current machine learning models, especially with limited data? (a) Speed of computation (b) Ability to generalize from few examples (c) Processing

large volumes of structured data (d) Performing repetitive tasks accurately

40. 25. Identifying whether a given image contains a 'cat' or 'dog' based on a trained model is an example of which type of machine learning task? a) Unsupervised Clustering b) Supervised Classification c) Reinforcement Learning d) Unsupervised Anomaly Detection

41. 14. Briefly discuss the concept of "generalization" in Machine Learning and why it is important for a model's performance.

42. 50. The primary purpose of the scikit-learn library in Python is to provide tools for which of the following? a) Building web applications b) Numerical optimization c) Traditional machine learning algorithms (classification, regression, clustering) d) Natural language processing only

43. 59. Why is Python often considered the preferred programming language for machine learning by many developers? a) It is faster than C++ for complex calculations. b) It has a steep learning curve but offers unique functionalities. c) It has a rich ecosystem of specialized libraries, ease of use, and strong community support. d) It is primarily used for front-end web development.

44. 33. Fraud detection in financial services is primarily an application of which ML paradigm? (a) Supervised Learning (classification) (b) Unsupervised Learning (clustering) (c) Reinforcement Learning (d) Feature Engineering

45. 1. Which of the following best defines Machine Learning? (a) Programming computers to perform tasks by explicitly coding rules. (b) Enabling computers to learn from data without being explicitly programmed. (c) The study of algorithms for building complex software systems. (d) Creating intelligent robots that mimic human behavior.

46. 11. Explain why "data" is considered the most critical component in any Machine Learning project.

47. 44. Jupyter Notebooks are widely used in machine learning for: (a) Deploying production models (b) Creating interactive visualizations and prototyping code (c) Managing large-scale data storage (d) Performing high-performance distributed training

48. 30. While Supervised, Unsupervised, and Reinforcement Learning are the main types, some modern approaches combine elements. Which of the following refers to a paradigm that uses a small amount of labeled data with a large amount of unlabeled data during training? a) Active Learning b) Semi-supervised Learning c) Transfer Learning d) Federated Learning

49. 18. Which of the following tasks is a classic example of Reinforcement Learning? a) Predicting the likelihood of customer churn based on past behavior. b) Grouping similar news articles into categories without prior labels. c) A self-driving car learning to navigate a city by maximizing rewards for correct actions. d) Classifying emails as spam or not spam.

50. 34. A key difference between human learning and machine learning is that machine learning typically requires: (a) Extensive prior knowledge and context (b) A structured dataset for training (c) Emotional intelligence for decision making (d) Minimal computational resources

51. 54. In Supervised Machine Learning, what is the essential characteristic that distinguishes it from Unsupervised Machine Learning? a) It works with unlabeled data. b) It does not require any historical data. c) It learns from data that has corresponding output labels. d) It focuses on finding hidden patterns without predefined outcomes.

52. 24. Consider a scenario where an AI is being trained to play a complex board game like chess, learning by playing against itself and being rewarded for winning moves. This best describes an application of: a) Classification b) Regression c) Reinforcement Learning d) Anomaly Detection
53. 23. When data scientists aim to reduce the number of features in a dataset while retaining most of the important information, they are likely using an algorithm from which machine learning paradigm? a) Supervised Learning b) Unsupervised Learning (specifically Dimensionality Reduction) c) Reinforcement Learning d) Semi-supervised Learning
54. 43. Which cloud platform offers services like Amazon SageMaker for building, training, and deploying machine learning models? (a) Microsoft Azure (b) Google Cloud Platform (c) Amazon Web Services (AWS) (d) IBM Cloud
55. 13. What is the purpose of a "training dataset" and a "test dataset" in Machine Learning?
56. 60. Which deep learning framework, developed by Facebook's AI Research lab, is known for its dynamic computation graph and is popular among researchers for its flexibility? a) TensorFlow b) Keras c) PyTorch d) MXNet
57. 9. Describe two key similarities between how humans learn and how machines are designed to learn.
58. 53. The process of transforming raw data into features that better represent the underlying problem to the predictive models, thereby improving model accuracy, is known as: a) Data Normalization b) Model Evaluation c) Feature Engineering d) Data Augmentation
59. 17. In Unsupervised Learning, what is the primary characteristic of the training data? a) It consists of input data paired with desired output labels. b) It consists of input data without any corresponding output labels. c) It is generated through an agent's interaction with an environment. d) It is a small set of labeled data combined with a large set of unlabeled data.
60. 28. Which type of machine learning is most suitable for tasks like anomaly detection, where the goal is to identify rare occurrences that deviate significantly from the majority of the data without prior examples of anomalies? a) Supervised Classification b) Reinforcement Learning c) Unsupervised Learning d) Semi-supervised Learning

ANSWER KEY

1. (c)
2. (b)
3. (c)
4. (b)
5. (c)
6. (c)
7. (c)
8. (b)
9. (c)
10. (d)
11. (c)
12. (c)
13. (c)
14. (c)
15. (c)
16. (b)
17. (b)
18. (c)
19. (c)
20. (b)
21. (c)
22. (c)
23. (c)
24. (c)
25. Spam detection; Recommendation systems.
26. (b)
27. (b)

28. (b)

29. (c)

30. Machine Learning is a field of AI enabling computers to learn from data without explicit programming. Its objective is to build systems that automatically learn patterns, make predictions, or take actions based on data.

31. (b)

32. (c)

33. (a) Supervised Learning (b) Reinforcement Learning (c) Unsupervised Learning

34. (c)

35. (b)

36. (d)

37. (b)

38. TensorFlow: An open-source library for numerical computation and large-scale machine learning, widely used for building and training deep neural networks.

39. (b)

40. (b)

41. Generalization is a model's ability to perform well on new, unseen data. It's important for ensuring the model is applicable and accurate in real-world scenarios beyond its training data.

42. (c)

43. (c)

44. (a)

45. (b)

46. Data is the "experience" from which a model learns; without quality data, a model cannot effectively learn patterns or make accurate predictions.

47. (b)

48. (b)

49. (c)

50. (b)

51. (c)

52. (c)

53. (b)

54. (c)

55. Training dataset: Used to teach the model to learn patterns. Test dataset: Used to evaluate the trained model's performance on unseen data.

56. (c)

57. Both learn from experience/data; both identify patterns.

58. (c)

59. (b)

60. (c)